

# Neural Network Framework Implementation

## A Step-by-Step Guide to Building Your Own Deep Learning Library

Morteza Maleki

Tarbiat Modares University

May 21, 2025

# Table of Contents

- 1 Introduction to the Framework
- 2 Backpropagation Fundamentals
- 3 Activation Functions
- 4 Initializers
- 5 Layers
- 6 Loss Functions
- 7 Models
- 8 Optimizers
- 9 Putting It All Together

# Overview of the NPDL Framework

- A NumPy-based Deep Learning Framework
- Modular design with separate components:
  - Activation functions
  - Weight initializers
  - Neural network layers
  - Loss functions
  - Model construction
  - Optimizers
- Implementation follows object-oriented principles
- Designed for educational purposes

# Backpropagation - The Core of Neural Networks

- Backpropagation: The algorithm that powers neural network learning
- Key components of a neural network training loop:
  - ① **Forward pass:** Compute predictions
  - ② **Loss calculation:** Measure error
  - ③ **Backward pass:** Compute gradients
  - ④ **Parameter update:** Improve model
- Focus of this section: Understanding how gradients flow backward through the network
- Goal: Connect mathematical theory to our code implementation

# The Chain Rule - Foundation of Backpropagation

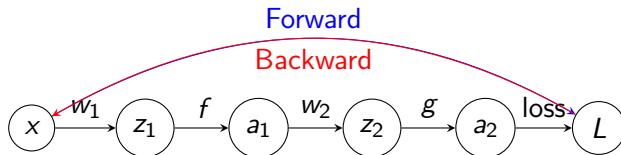
- Neural networks are composed of nested functions
- Chain rule from calculus: If  $y = f(g(x))$ , then  $\frac{dy}{dx} = \frac{dy}{dg} \cdot \frac{dg}{dx}$
- For neural networks with many layers:

$$\frac{\partial L}{\partial w_{ij}} = \frac{\partial L}{\partial y_k} \cdot \frac{\partial y_k}{\partial a_j} \cdot \frac{\partial a_j}{\partial z_j} \cdot \frac{\partial z_j}{\partial w_{ij}}$$

Where:

- $L$  is the loss
- $y_k$  is the output
- $a_j$  is the activation
- $z_j$  is the pre-activation
- $w_{ij}$  is a weight parameter

# Computational Graph Perspective



- Forward pass (blue): Compute outputs from inputs
- Backward pass (red): Propagate gradients from outputs to inputs
- Each node knows:
  - How to compute its output (forward)
  - How to compute gradients w.r.t. its inputs (backward)

# A 2-Layer Neural Network Example

Consider a simple 2-layer neural network:

$$z^{[1]} = W^{[1]}x + b^{[1]} \quad (1)$$

$$a^{[1]} = f(z^{[1]}) \quad (2)$$

$$z^{[2]} = W^{[2]}a^{[1]} + b^{[2]} \quad (3)$$

$$\hat{y} = g(z^{[2]}) \quad (4)$$

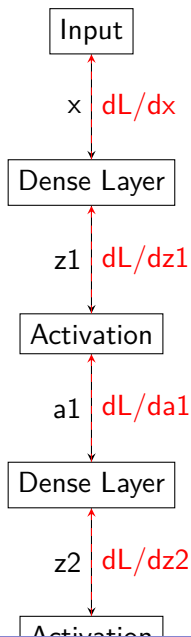
$$L = \text{loss}(\hat{y}, y) \quad (5)$$

- Computing gradients using the chain rule:

$$\frac{\partial L}{\partial W^{[2]}} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z^{[2]}} \cdot \frac{\partial z^{[2]}}{\partial W^{[2]}} \quad (6)$$

$$\frac{\partial L}{\partial W^{[1]}} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z^{[2]}} \cdot \frac{\partial z^{[2]}}{\partial a^{[1]}} \cdot \frac{\partial a^{[1]}}{\partial z^{[1]}} \cdot \frac{\partial z^{[1]}}{\partial W^{[1]}} \quad (7)$$

# Gradient Flow in Our Framework





# The Layer Interface in Our Framework

- Each layer in our framework follows the same interface:
  - 1 `forward()` - Computes layer output
  - 2 `backward()` - Computes gradients
  - 3 Stores parameters in `self.params`
  - 4 Stores gradients in `self.grads`
- This consistent interface enables automatic differentiation
- Models chain these layers together in both passes:
  - Forward: from first layer to last
  - Backward: from last layer to first

## Example: Gradient Flow in a Dense Layer

```
1 class Dense(Layer):
2     def forward(self, inputs):
3         """Forward pass:  $y = Wx + b$ """
4         self.inputs = inputs # Store for backward pass
5         return np.dot(inputs, self.weights) + self.bias
6
7     def backward(self, output_grad):
8         """Backward pass: Compute gradients"""
9         # Compute weight gradients (L/W)
10        self.grads['weights'] = np.dot(self.inputs.T, output_grad)
11
12        # Compute bias gradients (L/b)
13        self.grads['bias'] = np.sum(output_grad, axis=0, keepdims=True)
14
15        # Compute input gradients (L/x) to pass to previous layer
16        return np.dot(output_grad, self.weights.T)
17
```

- Dense layer implements matrix multiplication:  $y = Wx + b$
- Backward pass uses calculus rules:

- $\frac{\partial L}{\partial W} = x^T \cdot \frac{\partial L}{\partial y}$
- $\frac{\partial L}{\partial b} = \sum \frac{\partial L}{\partial y}$
- $\frac{\partial L}{\partial x} = \frac{\partial L}{\partial y} \cdot W^T$

# Example: Gradient Flow in an Activation Function

```
1 class ReLU(Activation):
2     def forward(self, x):
3         """Forward pass: f(x) = max(0, x)"""
4         self.input = x # Store for backward pass
5         self.output = np.maximum(0, x)
6         return self.output
7
8     def backward(self, output_grad):
9         """Backward pass: Compute gradients"""
10        # Derivative of ReLU: 1 if x > 0, else 0
11        relu_grad = (self.input > 0).astype(float)
12
13        # Chain rule: multiply upstream gradient with local gradient
14        return output_grad * relu_grad
15
```

- Activation functions apply element-wise operations
- Backward pass computes local derivatives and applies chain rule

- For ReLU:  $\frac{d}{dx} \max(0, x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$

# How Backward Passes Connect in the Framework

- Sequential model chains the backward passes automatically
- The flow of gradients:
  - 1 Loss function computes initial gradient:  $\frac{\partial L}{\partial \hat{y}}$
  - 2 Each layer receives gradient w.r.t. its output:  $\frac{\partial L}{\partial \text{out}}$
  - 3 Each layer computes:
    - Gradients for its parameters:  $\frac{\partial L}{\partial \theta}$
    - Gradient w.r.t. its input:  $\frac{\partial L}{\partial \text{in}}$
  - 4 Gradient w.r.t. input gets passed to previous layer
- Key insight: Each layer only needs to compute local gradients
- The framework handles chaining the gradients together

# Sequential Model's Backward Pass

```
1 class Sequential(object):
2     def __init__(self, layers=None):
3         self.layers = layers if layers is not None else []
4
5     def forward(self, inputs):
6         """Forward pass through all layers in sequence"""
7         for layer in self.layers:
8             inputs = layer.forward(inputs)
9         return inputs
10
11    def backward(self, grad):
12        """Backward pass through all layers in reverse"""
13        for layer in reversed(self.layers):
14            grad = layer.backward(grad)
15        return grad
16
```

- Sequential model chains layers together
- **Forward:** Process inputs through layers in order
- **Backward:** Process gradients through layers in reverse order
- Each layer receives the gradient from the layer ahead
- The full chain rule is automatically implemented

# Implementation to Mathematics Mapping

Component	Forward	Backward
Dense Layer	$z = Wx + b$	$\begin{aligned}\frac{\partial L}{\partial W} &= x^T \frac{\partial L}{\partial z} \\ \frac{\partial L}{\partial b} &= \sum \frac{\partial L}{\partial z} \\ \frac{\partial L}{\partial x} &= \frac{\partial L}{\partial z} W^T\end{aligned}$
ReLU	$a = \max(0, z)$	$\frac{\partial L}{\partial z} = \frac{\partial L}{\partial a} \cdot \mathbf{1}_{z>0}$
Sigmoid	$a = \frac{1}{1+e^{-z}}$	$\frac{\partial L}{\partial z} = \frac{\partial L}{\partial a} \cdot a(1-a)$
Softmax	$a_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$	Complex Jacobian matrix
MSE Loss	$L = \frac{1}{2}(y - \hat{y})^2$	$\frac{\partial L}{\partial \hat{y}} = \hat{y} - y$

- Each component has a mathematical operation and its derivative
- The code implementation directly follows these equations
- Understanding the math makes the code implementation clear

# Updating Parameters with Optimizers

- After computing gradients, optimizers update parameters
- Stochastic Gradient Descent (SGD):

$$\theta_{t+1} = \theta_t - \alpha \nabla_{\theta} J(\theta)$$

- SGD with momentum:

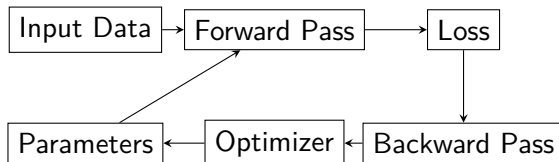
$$v_{t+1} = \gamma v_t + \alpha \nabla_{\theta} J(\theta) \quad (8)$$

$$\theta_{t+1} = \theta_t - v_{t+1} \quad (9)$$

- Each optimizer implements the `update(layer)` method:
  - Reads gradients from `layer.grads`
  - Updates parameters in `layer.params`
  - May maintain its own state (e.g., momentum)

# The Complete Training Loop

- ➊ **Forward pass:** Input  $\rightarrow$  Model  $\rightarrow$  Prediction
- ➋ **Loss calculation:** Compare prediction with target
- ➌ **Backward pass:**
  - Start with loss gradient
  - Propagate through model in reverse
  - Collect parameter gradients
- ➍ **Parameter update:**
  - Apply optimizer to update weights using gradients
- ➎ Repeat for many examples and epochs





# Summary: Backpropagation in Our Framework

- Key components implementing backpropagation:
  - **Layers:** Compute outputs and gradients
  - **Activations:** Add non-linearity and corresponding gradients
  - **Loss functions:** Measure error and provide initial gradient
  - **Optimizers:** Update parameters using gradients
  - **Models:** Chain components together
- Design principles:
  - Modularity: Each component has specific responsibility
  - Unified interface: `forward()`/`backward()` methods
  - Mathematical correspondence: Code follows math directly
  - Automatic differentiation: Chain rule applied automatically
- Next, we'll explore each component in detail

# Activation Functions - Overview

- Activation functions introduce non-linearity
- Base abstract class with common interface
- Implementations of common activation functions:
  - Sigmoid
  - Tanh
  - ReLU
  - Softmax
- Each has forward and backward methods

# Activation Functions - Base Class

```
1 class Activation(object):
2     """Base class for all activation functions."""
3
4     def forward(self, x):
5         """Forward pass."""
6         raise NotImplementedError
7
8     def backward(self, output_grad):
9         """Backward pass."""
10        raise NotImplementedError
```

# Activation Functions - Sigmoid

```
1 class Sigmoid(Activation):
2     """Sigmoid activation function."""
3
4     def forward(self, x):
5         """Forward pass for sigmoid function.
6
7         Args:
8             x: Input numpy array.
9         """
10        self.output = 1.0 / (1.0 + np.exp(-x))
11        return self.output
12
13    def backward(self, output_grad):
14        """Backward pass for sigmoid function.
15
16        Args:
17            output_grad: Gradient of the cost with respect to the
18            output.
19        """
20        return output_grad * self.output * (1 - self.output)
```

# Sigmoid Activation Function - Mathematical Definition

- The sigmoid function is defined as:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- Key properties:
  - Outputs values between 0 and 1
  - Smooth and differentiable everywhere
  - Historically popular but prone to vanishing gradient
  - S-shaped curve (sigmoid shape)
- Used primarily in:
  - Binary classification (output layer)
  - Early neural networks (mostly replaced by ReLU)
  - Gates in recurrent neural networks (LSTM, GRU)

# Sigmoid Activation Function - Gradient Derivation

$$\sigma(x) = \frac{1}{1 + e^{-x}} \quad (10)$$

$$\frac{d\sigma(x)}{dx} = \frac{d}{dx} \left( \frac{1}{1 + e^{-x}} \right) \quad (11)$$

$$= \frac{d}{dx} (1 + e^{-x})^{-1} \quad (12)$$

$$= -(1 + e^{-x})^{-2} \cdot \frac{d}{dx} (1 + e^{-x}) \quad (13)$$

$$= -(1 + e^{-x})^{-2} \cdot (-e^{-x}) \quad (14)$$

$$= \frac{e^{-x}}{(1 + e^{-x})^2} \quad (15)$$

$$= \frac{1}{1 + e^{-x}} \cdot \frac{e^{-x}}{1 + e^{-x}} \quad (16)$$

$$= \frac{1}{1 + e^{-x}} \cdot \left( 1 - \frac{1}{1 + e^{-x}} \right) \quad (17)$$

# Sigmoid Activation Function - Implementation Details

- Forward pass:
  - Calculate  $\sigma(x) = \frac{1}{1+e^{-x}}$
  - Store output for use in backward pass
  - Handle numerical stability with clipping for extreme values
- Backward pass:
  - Use the elegant form of gradient:  $\sigma'(x) = \sigma(x) \cdot (1 - \sigma(x))$
  - Multiply input gradient by this derivative (chain rule)
  - No need to reference original input - only output is needed
- Implementation challenges:
  - Saturation for large positive/negative inputs
  - Vanishing gradient problem when chaining multiple sigmoids

# ReLU Activation Function - Detailed Overview

- Rectified Linear Unit defined as:

$$\text{ReLU}(x) = \max(0, x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$

- Key properties:
  - Simple, computationally efficient
  - Non-linear despite simple form
  - Sparse activation (many neurons output zero)
  - No vanishing gradient for positive inputs
  - Allows for deeper networks
- Limitations:
  - "Dying ReLU" problem when neurons get stuck at 0
  - Non-zero centered outputs
  - Unbounded positive activation



# ReLU Activation Function - Gradient Analysis

$$\frac{d\text{ReLU}(x)}{dx} = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x < 0 \\ \text{undefined} & \text{if } x = 0 \end{cases} \quad (19)$$

- Gradient is either 0 or 1 (easy to compute)
- No saturation for positive values (solves vanishing gradient)
- Promotes sparsity in the network
- Gradient at  $x = 0$  is technically undefined, but usually set to 0 or 1
- Computationally efficient gradient - just check if input was positive

# ReLU Variants - LeakyReLU, PReLU, ELU

- **LeakyReLU**: Allows small negative values with a fixed slope

$$\text{LeakyReLU}(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha x & \text{if } x \leq 0 \end{cases}$$

where  $\alpha$  is a small constant (e.g., 0.01)

- **Parametric ReLU (PReLU)**: Learns the slope parameter  $\alpha$  during training
- **Exponential Linear Unit (ELU)**:

$$\text{ELU}(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha(e^x - 1) & \text{if } x \leq 0 \end{cases}$$

- These variants help address the "dying ReLU" problem while maintaining ReLU's advantages

# Activation Functions - ReLU

```
1 class ReLU(Activation):
2     """ReLU activation function."""
3
4     def forward(self, x):
5         """Forward pass for ReLU function.
6
7         Args:
8             x: Input numpy array.
9         """
10        self.input = x
11        self.output = np.maximum(x, 0)
12        return self.output
13
14    def backward(self, output_grad):
15        """Backward pass for ReLU function.
16
17        Args:
18            output_grad: Gradient of the cost with respect to the
19            output.
20        """
21        return output_grad * (self.input > 0)
```

# Softmax Activation Function - Comprehensive Overview

- Softmax converts a vector of values into a probability distribution:

$$\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

- Key properties:
  - Outputs are in range  $[0,1]$  and sum to 1
  - Preserves relative order of inputs (monotonic)
  - Emphasizes largest values, suppresses smaller ones
  - Not element-wise (depends on all input values)
- Use cases:
  - Output layer for multi-class classification
  - Attention mechanisms in transformers
  - Any scenario requiring normalized probabilities

# Softmax Activation Function - Numerical Stability

- Naive implementation can cause numerical overflow:

$$\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

- Stabilized implementation:

$$\text{softmax}(x_i) = \frac{e^{x_i - \max(x)}}{\sum_{j=1}^n e^{x_j - \max(x)}}$$

- Subtracting the maximum value prevents overflow:
  - Doesn't change the output (same relative proportions)
  - Ensures at least one exponent equals 1 (when  $x_i = \max(x)$ )
  - Makes largest exponent manageable
  - Critical for stable computation with larger numbers

# Softmax Activation Function - Gradient

- The gradient is a Jacobian matrix:

$$\frac{\partial \text{softmax}(x_i)}{\partial x_j} = \begin{cases} \text{softmax}(x_i)(1 - \text{softmax}(x_i)) & \text{if } i = j \\ -\text{softmax}(x_i)\text{softmax}(x_j) & \text{if } i \neq j \end{cases}$$

- In matrix form:

$$\nabla_{\mathbf{x}} \text{softmax}(\mathbf{x}) = \text{diag}(\text{softmax}(\mathbf{x})) - \text{softmax}(\mathbf{x}) \otimes \text{softmax}(\mathbf{x})$$

- In practice, often combined with cross-entropy loss for simplification
- When used with cross-entropy, gradient simplifies to:  $(\hat{y} - y)$

# Activation Functions - Softmax

```
1 class Softmax(Activation):
2     """Softmax activation function."""
3
4     def forward(self, x):
5         """Forward pass for Softmax function.
6
7         Args:
8             x: Input numpy array.
9         """
10        exp_values = np.exp(x - np.max(x, axis=1, keepdims=True))
11        self.output = exp_values / np.sum(exp_values,
12                                           axis=1, keepdims=True)
13        return self.output
14
15    def backward(self, output_grad):
16        """Backward pass for Softmax function.
17
18        Args:
19            output_grad: Gradient of the cost with respect to the
20            output.
21        """
22        # Simplified backward pass
23        return output_grad
```

# Neural Network Layers - Overview

- Building blocks of neural networks
- Base Layer interface
- Implementations:
  - Dense (Fully connected)
  - Dropout (Regularization)
  - Flatten (Reshaping)
- Each layer implements forward and backward passes



# Layers - Base Class

```
1 class Layer(object):
2     """Base class for all layers."""
3
4     def __init__(self):
5         """Initialize the layer."""
6         self.params = {}
7         self.grads = {}
8
9     def forward(self, inputs):
10        """Forward pass.
11
12        Args:
13            inputs: Input data.
14        """
15        raise NotImplementedError
16
17    def backward(self, output_grad):
18        """Backward pass.
19
20        Args:
21            output_grad: Gradient of the cost with respect to the
22            output.
23        """
24        raise NotImplementedError
```

# Layers - Dense Layer (1/2)

```
1 class Dense(Layer):
2     """Fully connected layer."""
3
4     def __init__(self, n_units, input_shape=None,
5                   weight_initializer=None,
6                   bias_initializer=None):
7         """Initialize the dense layer.
8
9         Args:
10             n_units: Number of output units.
11             input_shape: Shape of the input data.
12             weight_initializer: Weight initializer.
13             bias_initializer: Bias initializer.
14         """
15         super(Dense, self).__init__()
16         self.n_units = n_units
17         self.input_shape = input_shape
18
19         self.weight_initializer = weight_initializer
20         if self.weight_initializer is None:
21             self.weight_initializer = HeNormal()
22
23         self.bias_initializer = bias_initializer
24         if self.bias_initializer is None:
25             self.bias_initializer = Zero()
```

## Layers - Dense Layer (2/2)

```
1  def forward(self, inputs):
2      """Forward pass for dense layer.
3
4      Args:
5          inputs: Input data.
6      """
7      self.inputs = inputs
8
9      if not hasattr(self, 'weights'):
10         self.weights = self.weight_initializer(
11             (self.input_shape, self.n_units))
12         self.bias = self.bias_initializer((1, self.n_units))
13
14         self.params['weights'] = self.weights
15         self.params['bias'] = self.bias
16
17     return np.dot(self.inputs, self.weights) + self.bias
18
19 def backward(self, output_grad):
20     """Backward pass for dense layer.
21
22     Args:
23         output_grad: Gradient of the cost with respect to the
24         output.
25     """
```

# Layers - Dropout Layer

```
1 class Dropout(Layer):
2     """Dropout layer."""
3
4     def __init__(self, dropout_rate):
5         """Initialize the dropout layer.
6
7         Args:
8             dropout_rate: Dropout rate.
9         """
10        super(Dropout, self).__init__()
11        self.dropout_rate = dropout_rate
12
13    def forward(self, inputs, training=True):
14        """Forward pass for dropout layer.
15
16        Args:
17            inputs: Input data.
18            training: Whether in training mode.
19        """
20        self.inputs = inputs
21
22        if training:
23            self.mask = np.random.binomial(
24                1, 1 - self.dropout_rate, size=inputs.shape) / (1 -
25                self.dropout_rate)
```

# Loss Functions - Overview

- Measure model performance
- Base Loss abstract class
- Common loss functions:
  - Mean Squared Error (MSE)
  - Categorical Cross-Entropy
  - Binary Cross-Entropy
- Each implements forward and gradient calculations

# Loss Functions - Base Class

```
1 class Loss(object):
2     """Base class for all loss functions."""
3
4     def forward(self, y_true, y_pred):
5         """Forward pass.
6
7         Args:
8             y_true: Ground truth values.
9             y_pred: Predicted values.
10        """
11        raise NotImplementedError
12
13    def gradient(self, y_true, y_pred):
14        """Gradient of the loss function.
15
16        Args:
17            y_true: Ground truth values.
18            y_pred: Predicted values.
19        """
20        raise NotImplementedError
```

# Loss Functions - MSE

```
1 class MeanSquaredError(Loss):
2     """Mean squared error loss function."""
3
4     def forward(self, y_true, y_pred):
5         """Forward pass for mean squared error.
6
7         Args:
8             y_true: Ground truth values.
9             y_pred: Predicted values.
10        """
11        return 0.5 * np.power(y_pred - y_true, 2).mean()
12
13    def gradient(self, y_true, y_pred):
14        """Gradient of the mean squared error.
15
16        Args:
17            y_true: Ground truth values.
18            y_pred: Predicted values.
19        """
20        return y_pred - y_true
```

# Loss Functions - Categorical Cross-Entropy

```
1 class CategoricalCrossentropy(Loss):
2     """Categorical cross-entropy loss function."""
3
4     def forward(self, y_true, y_pred):
5         """Forward pass for categorical cross-entropy.
6
7         Args:
8             y_true: Ground truth values.
9             y_pred: Predicted values.
10        """
11        # Clip to avoid log(0)
12        y_pred = np.clip(y_pred, 1e-15, 1 - 1e-15)
13        return -np.sum(y_true * np.log(y_pred)) / y_true.shape[0]
14
15    def gradient(self, y_true, y_pred):
16        """Gradient of the categorical cross-entropy.
17
18        Args:
19            y_true: Ground truth values.
20            y_pred: Predicted values.
21        """
22        # Clip to avoid division by zero
23        y_pred = np.clip(y_pred, 1e-15, 1 - 1e-15)
24        return -y_true / y_pred / y_true.shape[0]
```



# Models - Overview

- Sequential model for chaining layers
- Methods for training and evaluation
- Forward/backward pass implementation
- Training loop with batch processing
- Model evaluation and prediction

# Models - Sequential (1/2)

```
1 class Sequential(object):
2     """Sequential model."""
3
4     def __init__(self, layers=None):
5         """Initialize the model.
6
7         Args:
8             layers: List of layers.
9         """
10        self.layers = layers if layers is not None else []
11
12    def add(self, layer):
13        """Add a layer to the model.
14
15        Args:
16            layer: Layer to add.
17        """
18        self.layers.append(layer)
```

# Models - Sequential (2/2)

```
1  def forward(self, inputs, training=True):
2      """Forward pass.
3
4      Args:
5          inputs: Input data.
6          training: Whether in training mode.
7      """
8      for layer in self.layers:
9          if hasattr(layer, 'training'):
10              inputs = layer.forward(inputs, training)
11          else:
12              inputs = layer.forward(inputs)
13      return inputs
14
15  def backward(self, grad):
16      """Backward pass.
17
18      Args:
19          grad: Gradient of the cost with respect to the output.
20      """
21      for layer in reversed(self.layers):
22          grad = layer.backward(grad)
23      return grad
```

# Models - Training

```
1  def fit(self, x, y, epochs=100, batch_size=32,
2      loss_fn=None, optimizer=None,
3      validation_data=None, verbose=True):
4      """Train the model.
5
6      Args:
7          x: Input data.
8          y: Target data.
9          epochs: Number of epochs.
10         batch_size: Batch size.
11         loss_fn: Loss function.
12         optimizer: Optimizer.
13         validation_data: Validation data.
14         verbose: Whether to print progress.
15     """
16     if loss_fn is None:
17         loss_fn = MeanSquaredError()
18
19     if optimizer is None:
20         optimizer = SGD()
21
22     # Training loop implementation
23     for epoch in range(epochs):
24         # Process mini-batches
25         # Update weights using optimizer
```

# Optimizers - Overview

- Update model parameters based on gradients
- Base Optimizer abstract class
- Common optimization algorithms:
  - SGD (Stochastic Gradient Descent)
  - Adam (Adaptive Moment Estimation)
  - RMSprop
- Each implements the update method

# Optimizers - Base Class

```
1 class Optimizer(object):
2     """Base class for all optimizers."""
3
4     def __init__(self, learning_rate=0.01):
5         """Initialize the optimizer.
6
7         Args:
8             learning_rate: Learning rate.
9         """
10        self.learning_rate = learning_rate
11
12    def update(self, layer):
13        """Update the layer weights.
14
15        Args:
16            layer: Layer to update.
17        """
18        raise NotImplementedError
```

# Optimizers - SGD

```
1 class SGD(Optimizer):
2     """Stochastic gradient descent optimizer."""
3
4     def __init__(self, learning_rate=0.01, momentum=0.0):
5         """Initialize the SGD optimizer.
6
7         Args:
8             learning_rate: Learning rate.
9             momentum: Momentum factor.
10        """
11        super(SGD, self).__init__(learning_rate)
12        self.momentum = momentum
13        self.velocity = {}
14
15    def update(self, layer):
16        """Update the layer weights.
17
18        Args:
19            layer: Layer to update.
20        """
21        for param_name in layer.params:
22            # Initialize velocity for the parameter if not exists
23            if param_name not in self.velocity:
24                self.velocity[param_name] = np.zeros_like(
25                    layer.params[param_name])
```

# Optimizers - Adam

```
1 class Adam(Optimizer):
2     """Adam optimizer."""
3
4     def __init__(self, learning_rate=0.001, beta_1=0.9,
5                 beta_2=0.999, epsilon=1e-8):
6         """Initialize the Adam optimizer.
7
8         Args:
9             learning_rate: Learning rate.
10            beta_1: Exponential decay rate for first moment.
11            beta_2: Exponential decay rate for second moment.
12            epsilon: Small constant for numerical stability.
13        """
14        super(Adam, self).__init__(learning_rate)
15        self.beta_1 = beta_1
16        self.beta_2 = beta_2
17        self.epsilon = epsilon
18        self.m = {} # First moment
19        self.v = {} # Second moment
20        self.t = 0 # Timestep
21
22    def update(self, layer):
23        """Update implementation with moment calculations"""
```



# Building a Complete Neural Network

```
1 # Import all components
2 from npdl.models import Sequential
3 from npdl.layers import Dense, Dropout
4 from npdl.activations import ReLU, Softmax
5 from npdl.initializers import HeNormal, Zero
6 from npdl.losses import CategoricalCrossentropy
7 from npdl.optimizers import Adam
8
9 # Create a model
10 model = Sequential()
11 model.add(Dense(128, input_shape=784,
12                 weight_initializer=HeNormal(),
13                 bias_initializer=Zero()))
14 model.add(ReLU())
15 model.add(Dropout(0.2))
16 model.add(Dense(64))
17 model.add(ReLU())
18 model.add(Dropout(0.2))
19 model.add(Dense(10))
20 model.add(Softmax())
21
22 # Compile and train
23 model.fit(x_train, y_train, epochs=10, batch_size=32,
24         loss_fn=CategoricalCrossentropy(),
25         optimizer=Adam(learning_rate=0.001)).
```

# Summary and Next Steps

- We've covered the complete implementation of:
  - Activation functions
  - Weight initializers
  - Neural network layers
  - Loss functions
  - Model construction
  - Optimizers
- Possible extensions:
  - Convolutional layers
  - Recurrent layers
  - Batch normalization
  - More advanced optimizers
- Practical exercises to implement and test